

Deep Reinforcement learning

ML course – fall 1401

Dr Rezazadegan

Kiana Asgari

Zahra Janalizadeh




The Wonder of Reinforcement Learning



Coast Runners Game, played by RL agent.




Reinforcement Learning (RL): Key Concepts



Agent: takes actions, for example, a drone. The algorithm is the agent, and in real life we are the agents.



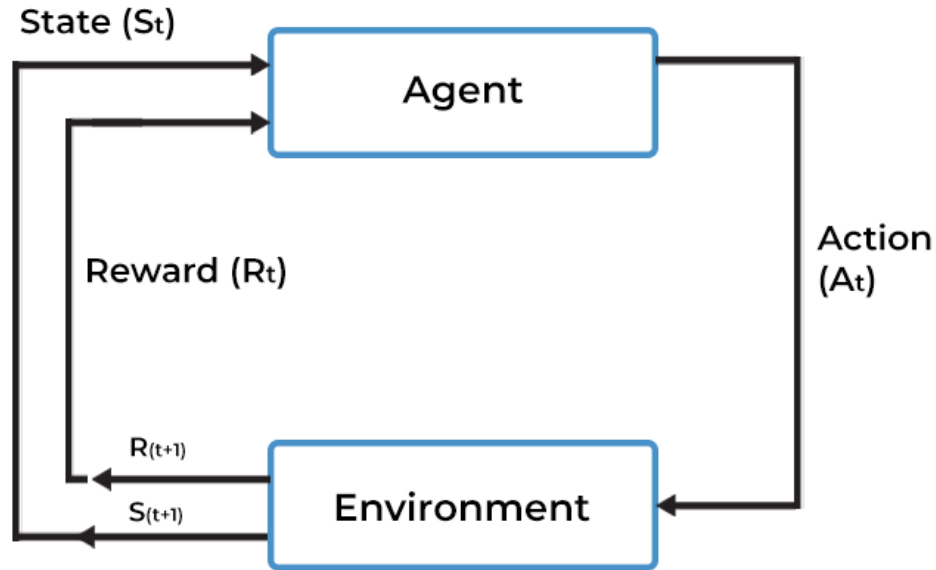
Reinforcement Learning (RL): Key Concepts



Environment: the world in which the agent exists and operates.




Reinforcement Learning (RL): Key Concepts



Action: a move the agent can make in the environment.




Reinforcement Learning (RL): Key Concepts



State: a situation which the agent perceives.



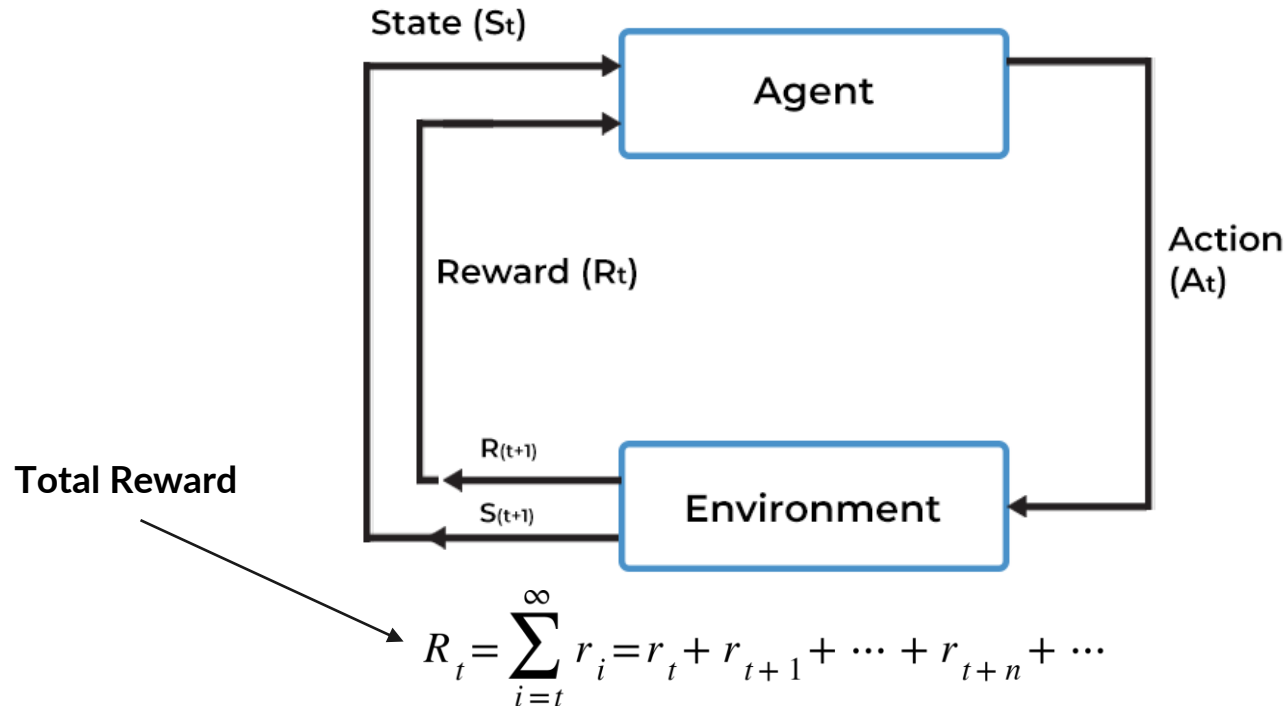
Reinforcement Learning (RL): Key Concepts



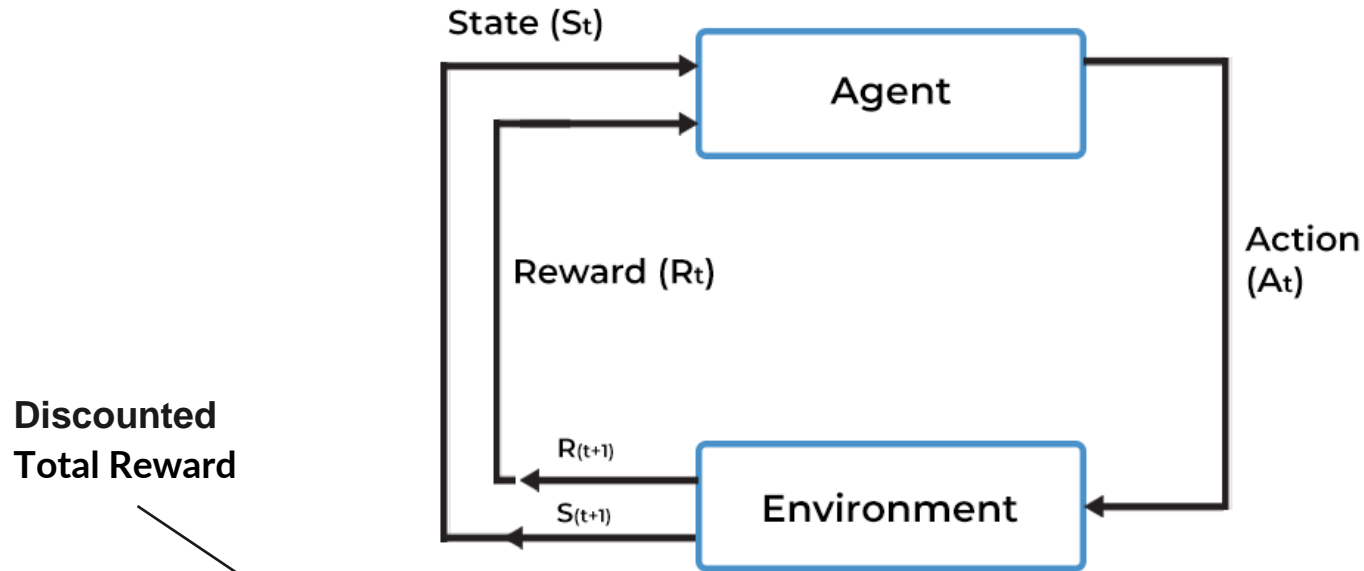
Reward: feedback that measures the success or failure of the agent's action.



Reinforcement Learning (RL): Key Concepts



Reinforcement Learning (RL): Key Concepts



Discounted
Total Reward

$$R_t = \sum_{i=t}^{\infty} \gamma^i r_i = \gamma^t r_t + \gamma^{t+1} r_{t+1} + \dots + \gamma^{t+n} r_{t+n} + \dots$$



Defining the Q-function

$$R_t = r_t + \gamma r_{t+1} + \gamma^2 r_{t+2} + \dots$$

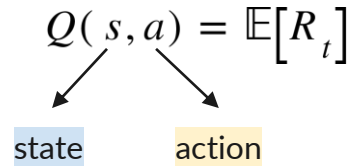
Total Reward, R_t , is the discounted sum of all rewards obtained from time t

$$Q(s, a) = \mathbb{E}[R_t]$$

The Q-function captures the **Expected Total Future Reward** an agent in state, s , can receive by executing a certain action, a



How to take actions given a Q-function?

$$Q(s, a) = \mathbb{E}[R_t]$$


state action

Ultimately, the agent needs a **policy** $\pi^*(s)$, to infer the best action to take at its state, s

Strategy: the policy should choose an action that maximizes the future reward

$$\pi^*(s) = \underset{a}{\operatorname{argmax}} Q(s, a)$$



Deep Reinforcement Learning Breakthrough



Deep Reinforcement Learning Breakthrough

→ DRL algorithms are able to handle high-dimensional state space!



Deep Reinforcement Learning Breakthrough

- DRL algorithms are able to handle high-dimensional state space!
- Beginning around 2013, **DeepMind** showed impressive learning results using deep RL to play **Atari** video games.



Deep Reinforcement Learning Breakthrough

- DRL algorithms are able to handle high-dimensional state space!
- Beginning around 2013, **DeepMind** showed impressive learning results using deep RL to play **Atari** video games.
- In 2016, DeepMind was able to solve **Go**, the most challenging classical problem for AI.





AlphaGo

AlphaGo! Making History

AlphaGo is the first computer program to defeat a professional human Go player, the first to defeat a Go world champion, and is arguably the strongest Go player in history.



Deep Reinforcement Learning Algorithms

Value Based


- Find $Q(s,a)$
 $a = \underset{a}{\operatorname{argmax}} Q(s,a)$

Policy Based

- Find $\pi(s)$
Sample $a \sim \pi(s)$




Deep Q Network (DQN)



Deep Q Learning




Deep Q Network (DQN)



Step 1: Choose your network structure and initial the DQN.

Deep Q Network (DQN)




Step 2: Choose an action from epsilon-greedy algorithm

$$\begin{cases} a = \underset{a}{\operatorname{argmax}} Q(s, a) & \text{with probability } 1 - \varepsilon \\ a = \text{random action} & \text{with probability } \varepsilon \end{cases}$$



Deep Q Network (DQN)




Step 3: Define your loss.

$$\text{LossFunction} = \overset{\text{TD-target is UNKNOWN!}}{\boxed{Q_{best}(s_t, a_t)}} - \overset{\text{Current Q-value}}{\boxed{Q(s_t, a_t)}}$$



Deep Q Network (DQN)




Step 3.5: Update your network weights using the Bellman Equation.

$$\text{LossFunction} = \underbrace{R_{t+1}}_{\text{Reward}} + \underbrace{\gamma}_{\text{Discount factor}} \underbrace{\max Q(s_{t+1}, a)}_{\text{Maximum next-state Q-value}} - \underbrace{Q(s_t, a_t)}_{\text{Current Q-value}}$$

Estimated TD-target





Downsides of Q-learning

Complexity:

- Can model scenarios where the action space is discrete and small
- Cannot handle continuous action spaces



Downsides of Q-learning

Complexity:

- Can model scenarios where the action space is discrete and small
- Cannot handle continuous action spaces

Imagine you want to predict steering wheel angle of a car!



Downsides of Q-learning

Complexity:

- Can model scenarios where the action space is discrete and small
- Cannot handle continuous action spaces

Imagine you want to predict steering wheel angle of a car!

Flexibility:

- Cannot learn stochastic policies, since policy is deterministically computed from the Q function



Downsides of Q-learning

Complexity:

- Can model scenarios where the action space is discrete and small
- Cannot handle continuous action spaces

Imagine you want to predict steering wheel angle of a car!

Flexibility:

- Cannot learn stochastic policies, since policy is deterministically computed from the Q function

What can we do to overcome this?



Downsides of Q-learning

Complexity:

- Can model scenarios where the action space is discrete and small
- Cannot handle continuous action spaces

Imagine you want to predict steering wheel angle of a car!

Flexibility:

- Cannot learn stochastic policies, since policy is deterministically computed from the Q function

What can we do to overcome this?

consider a new class of RL training algorithms: **Policy Gradient methods**



Policy Gradient (PG)

DQN (before)

Approximating Q and inferring
the optimal policy.

Policy Gradient

Directly optimize the policy!




Policy Gradient (PG)

DQN (before)

Approximating Q and inferring
the optimal policy.

Policy Gradient

Directly optimize the policy!



$$\sum_{a_i \in A} \pi(a_i|s) = 1$$

$$\pi(a|s) = P(\text{action}|\text{state})$$

Policy Gradient (PG): Training

1. Run a policy for a while
2. Increase probability of actions that lead to high rewards
3. Decrease probability of actions that lead to low/no rewards.



Policy Gradient (PG): Training

1. Run a policy for a while
2. Increase probability of actions that lead to high rewards
3. Decrease probability of actions that lead to low/no rewards.

```
function REINFORCE
  Initialize  $\theta$ 
  for  $episode \sim \pi_\theta$ 
     $\{s_i, a_i, r_i\}_{i=1}^{T-1} \leftarrow episode$ 
    for  $t = 1$  to  $T-1$ 
       $\nabla \leftarrow \nabla_\theta \log \pi_\theta(a_t | s_t) R_t$ 
       $\theta \leftarrow \theta + \alpha \nabla$ 
  return  $\theta$ 
```



Policy Gradient (PG): Training

1. Run a policy for a while
2. Increase probability of actions that lead to high rewards
3. Decrease probability of actions that lead to low/no rewards.

```
function REINFORCE
  Initialize  $\theta$ 
  for  $episode \sim \pi_\theta$ 
     $\{s_i, a_i, r_i\}_{i=1}^{T-1} \leftarrow episode$ 
    for  $t = 1$  to  $T-1$ 
       $\nabla \leftarrow \nabla_\theta \log \pi_\theta(a_t | s_t) R_t$ 
       $\theta \leftarrow \theta + \alpha \nabla$ 
  return  $\theta$ 
```

$$\nabla_\theta \log \pi_\theta(a_t | s_t) R_t$$

Log-likelihood of action

Reward



Stable Baseline3:

A reliable user-friendly implementation of DRL!



**Thank you for your
Attention! ^.^**

